

Problem Statement

NIDDK (National Institute of Diabetes and Digestive and Kidney Diseases) research creates knowledge about and treatments for the most chronic, costly, and consequential diseases. The dataset used in this project is originally from NIDDK. The objective is to predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Build a model to accurately predict whether the patients in the dataset have diabetes or not.

Dataset Description

The datasets consists of several medical predictor variables and one target variable (Outcome). Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and more.

Variables Description

- Pregnancies-Number of times pregnant
- Glucose- Plasma glucose concentration in an oral glucose tolerance test
- BloodPressure-Diastolic blood pressure (mm Hg)
- SkinThickness- Triceps skinfold thickness (mm)
- Insulin-Two hour serum insulin
- BMI-Body Mass Index
- DiabetesPedigreeFunction-Diabetes pedigree function
- Age-Age in years
- Outcome-Class variable (either 0 or 1). 268 of 768 values are 1, and the others are 0

Importing Necessary Libraries

In [190...

```
#Classic, Data Manipulation

import pandas as pd
import numpy as np

#Plots
import matplotlib.pyplot as plt
import seaborn as sns

#Data processing, metrics and modeling

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, p
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# to display Image files
from PIL import Image as PILImage
```

```
#ignore warning messages
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('C:/Users/vipul/Downloads/Project_2/Project 2/Healthcare - Diabetes
```

```
In [3]: df.head()
```

Out[3]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

```
In [4]: df.shape
```

Out[4]: (768, 9)

```
In [5]: df.describe()
```

Out[5]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

Unique

```
In [6]: for i in df.columns:
        print(i,df[i].unique(),'\n')
```

```
Pregnancies [ 6  1  8  0  5  3 10  2  4  7  9 11 13 15 17 12 14]

Glucose [148  85 183  89 137 116  78 115 197 125 110 168 139 189 166 100 118 107
103 126  99 196 119 143 147  97 145 117 109 158  88  92 122 138 102  90
111 180 133 106 171 159 146  71 105 101 176 150  73 187  84  44 141 114
 95 129  79   0  62 131 112 113  74  83 136  80 123  81 134 142 144  93
163 151  96 155  76 160 124 162 132 120 173 170 128 108 154  57 156 153
188 152 104  87  75 179 130 194 181 135 184 140 177 164  91 165  86 193
191 161 167  77 182 157 178  61  98 127  82  72 172  94 175 195  68 186
198 121  67 174 199  56 169 149  65 190]
```

BloodPressure [72 66 64 40 74 50 0 70 96 92 80 60 84 30 88 90 94 76

82 75 58 78 68 110 56 62 85 86 48 44 65 108 55 122 54 52
98 104 95 46 102 100 61 24 38 106 114]

SkinThickness [35 29 0 23 32 45 19 47 38 30 41 33 26 15 36 11 31 37 42 25 18 24 39 27

21 34 10 60 13 20 22 28 54 40 51 56 14 17 50 44 12 46 16 7 52 43 48 8
49 63 99]

Insulin [0 94 168 88 543 846 175 230 83 96 235 146 115 140 110 245 54 192

207 70 240 82 36 23 300 342 304 142 128 38 100 90 270 71 125 176
48 64 228 76 220 40 152 18 135 495 37 51 99 145 225 49 50 92
325 63 284 119 204 155 485 53 114 105 285 156 78 130 55 58 160 210
318 44 190 280 87 271 129 120 478 56 32 744 370 45 194 680 402 258
375 150 67 57 116 278 122 545 75 74 182 360 215 184 42 132 148 180
205 85 231 29 68 52 255 171 73 108 43 167 249 293 66 465 89 158
84 72 59 81 196 415 275 165 579 310 61 474 170 277 60 14 95 237
191 328 250 480 265 193 79 86 326 188 106 65 166 274 77 126 330 600
185 25 41 272 321 144 15 183 91 46 440 159 540 200 335 387 22 291
392 178 127 510 16 112]

BMI [33.6 26.6 23.3 28.1 43.1 25.6 31. 35.3 30.5 0. 37.6 38. 27.1 30.1

25.8 30. 45.8 29.6 43.3 34.6 39.3 35.4 39.8 29. 36.6 31.1 39.4 23.2
22.2 34.1 36. 31.6 24.8 19.9 27.6 24. 33.2 32.9 38.2 37.1 34. 40.2
22.7 45.4 27.4 42. 29.7 28. 39.1 19.4 24.2 24.4 33.7 34.7 23. 37.7
46.8 40.5 41.5 25. 25.4 32.8 32.5 42.7 19.6 28.9 28.6 43.4 35.1 32.
24.7 32.6 43.2 22.4 29.3 24.6 48.8 32.4 38.5 26.5 19.1 46.7 23.8 33.9
20.4 28.7 49.7 39. 26.1 22.5 39.6 29.5 34.3 37.4 33.3 31.2 28.2 53.2
34.2 26.8 55. 42.9 34.5 27.9 38.3 21.1 33.8 30.8 36.9 39.5 27.3 21.9
40.6 47.9 50. 25.2 40.9 37.2 44.2 29.9 31.9 28.4 43.5 32.7 67.1 45.
34.9 27.7 35.9 22.6 33.1 30.4 52.3 24.3 22.9 34.8 30.9 40.1 23.9 37.5
35.5 42.8 42.6 41.8 35.8 37.8 28.8 23.6 35.7 36.7 45.2 44. 46.2 35.
43.6 44.1 18.4 29.2 25.9 32.1 36.3 40. 25.1 27.5 45.6 27.8 24.9 25.3
37.9 27. 26. 38.7 20.8 36.1 30.7 32.3 52.9 21. 39.7 25.5 26.2 19.3
38.1 23.5 45.5 23.1 39.9 36.8 21.8 41. 42.2 34.4 27.2 36.5 29.8 39.2
38.4 36.2 48.3 20. 22.3 45.7 23.7 22.1 42.1 42.4 18.2 26.4 45.3 37.
24.5 32.2 59.4 21.2 26.7 30.2 46.1 41.3 38.8 35.2 42.3 40.7 46.5 33.5
37.3 30.3 26.3 21.7 36.4 28.5 26.9 38.6 31.3 19.5 20.1 40.8 23.4 28.3
38.9 57.3 35.6 49.6 44.6 24.1 44.5 41.2 49.3 46.3]

DiabetesPedigreeFunction [0.627 0.351 0.672 0.167 2.288 0.201 0.248 0.134 0.158 0.23 2 0.191 0.537

1.441 0.398 0.587 0.484 0.551 0.254 0.183 0.529 0.704 0.388 0.451 0.263
0.205 0.257 0.487 0.245 0.337 0.546 0.851 0.267 0.188 0.512 0.966 0.42
0.665 0.503 1.39 0.271 0.696 0.235 0.721 0.294 1.893 0.564 0.586 0.344
0.305 0.491 0.526 0.342 0.467 0.718 0.962 1.781 0.173 0.304 0.27 0.699
0.258 0.203 0.855 0.845 0.334 0.189 0.867 0.411 0.583 0.231 0.396 0.14
0.391 0.37 0.307 0.102 0.767 0.237 0.227 0.698 0.178 0.324 0.153 0.165
0.443 0.261 0.277 0.761 0.255 0.13 0.323 0.356 0.325 1.222 0.179 0.262
0.283 0.93 0.801 0.207 0.287 0.336 0.247 0.199 0.543 0.192 0.588 0.539
0.22 0.654 0.223 0.759 0.26 0.404 0.186 0.278 0.496 0.452 0.403 0.741
0.361 1.114 0.457 0.647 0.088 0.597 0.532 0.703 0.159 0.268 0.286 0.318
0.272 0.572 0.096 1.4 0.218 0.085 0.399 0.432 1.189 0.687 0.137 0.637
0.833 0.229 0.817 0.204 0.368 0.743 0.722 0.256 0.709 0.471 0.495 0.18
0.542 0.773 0.678 0.719 0.382 0.319 0.19 0.956 0.084 0.725 0.299 0.244
0.745 0.615 1.321 0.64 0.142 0.374 0.383 0.578 0.136 0.395 0.187 0.905
0.15 0.874 0.236 0.787 0.407 0.605 0.151 0.289 0.355 0.29 0.375 0.164
0.431 0.742 0.514 0.464 1.224 1.072 0.805 0.209 0.666 0.101 0.198 0.652
2.329 0.089 0.645 0.238 0.394 0.293 0.479 0.686 0.831 0.582 0.446 0.402
1.318 0.329 1.213 0.427 0.282 0.143 0.38 0.284 0.249 0.926 0.557 0.092
0.655 1.353 0.612 0.2 0.226 0.997 0.933 1.101 0.078 0.24 1.136 0.128
0.422 0.251 0.677 0.296 0.454 0.744 0.881 0.28 0.259 0.619 0.808 0.34
0.434 0.757 0.613 0.692 0.52 0.412 0.84 0.839 0.156 0.215 0.326 1.391
0.875 0.313 0.433 0.626 1.127 0.315 0.345 0.129 0.527 0.197 0.731 0.148
0.123 0.127 0.122 1.476 0.166 0.932 0.343 0.893 0.331 0.472 0.673 0.389
0.485 0.349 0.279 0.346 0.252 0.243 0.58 0.559 0.302 0.569 0.378 0.385
0.499 0.306 0.234 2.137 1.731 0.545 0.225 0.816 0.528 0.509 1.021 0.821

```

0.947 1.268 0.221 0.66 0.239 0.949 0.444 0.463 0.803 1.6 0.944 0.196
0.241 0.161 0.135 0.376 1.191 0.702 0.674 1.076 0.534 1.095 0.554 0.624
0.219 0.507 0.561 0.421 0.516 0.264 0.328 0.233 0.108 1.138 0.147 0.727
0.435 0.497 0.23 0.955 2.42 0.658 0.33 0.51 0.285 0.415 0.381 0.832
0.498 0.212 0.364 1.001 0.46 0.733 0.416 0.705 1.022 0.269 0.6 0.571
0.607 0.17 0.21 0.126 0.711 0.466 0.162 0.419 0.63 0.365 0.536 1.159
0.629 0.292 0.145 1.144 0.174 0.547 0.163 0.738 0.314 0.968 0.409 0.297
0.525 0.154 0.771 0.107 0.493 0.717 0.917 0.501 1.251 0.735 0.804 0.661
0.549 0.825 0.423 1.034 0.16 0.341 0.68 0.591 0.3 0.121 0.502 0.401
0.601 0.748 0.338 0.43 0.892 0.813 0.693 0.575 0.371 0.206 0.417 1.154
0.925 0.175 1.699 0.682 0.194 0.4 0.1 1.258 0.482 0.138 0.593 0.878
0.157 1.282 0.141 0.246 1.698 1.461 0.347 0.362 0.393 0.144 0.732 0.115
0.465 0.649 0.871 0.149 0.695 0.303 0.61 0.73 0.447 0.455 0.133 0.155
1.162 1.292 0.182 1.394 0.217 0.631 0.88 0.614 0.332 0.366 0.181 0.828
0.335 0.856 0.886 0.439 0.253 0.598 0.904 0.483 0.565 0.118 0.177 0.176
0.295 0.441 0.352 0.826 0.97 0.595 0.317 0.265 0.646 0.426 0.56 0.515
0.453 0.785 0.734 1.174 0.488 0.358 1.096 0.408 1.182 0.222 1.057 0.766
0.171]

```

```

Age [50 31 32 21 33 30 26 29 53 54 34 57 59 51 27 41 43 22 38 60 28 45 35 46
56 37 48 40 25 24 58 42 44 39 36 23 61 69 62 55 65 47 52 66 49 63 67 72
81 64 70 68]

```

```
Outcome [1 0]
```

```
In [7]: df.columns
```

```
Out[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

What all features have zero values

```
In [8]: for i in df.columns[1:-1]:
        l = len(df[df[i]==0])
        if l>=1:
            print(i,'---- has total {} Zero values'.format(l))
        else:
            print(i,'---- has no zero values and is good to go')
```

```

Glucose ---- has total 5 Zero values
BloodPressure ---- has total 35 Zero values
SkinThickness ---- has total 227 Zero values
Insulin ---- has total 374 Zero values
BMI ---- has total 11 Zero values
DiabetesPedigreeFunction ---- has no zero values and is good to go
Age ---- has no zero values and is good to go

```

A person can not have zero values for Glucose, Bloodpressure, SkinThickness, Insulin, BMI and Diabetes Pedigress Function. All these zero values don't make any sense hence these are nothing but the missing values. So we'll treat them with missing values imputation techniques

```
In [9]: df_copy = df.copy(deep = True)
        df_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] = df_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']].replace(0, np.nan)
```

```
In [10]: df_copy.isnull().sum()
```

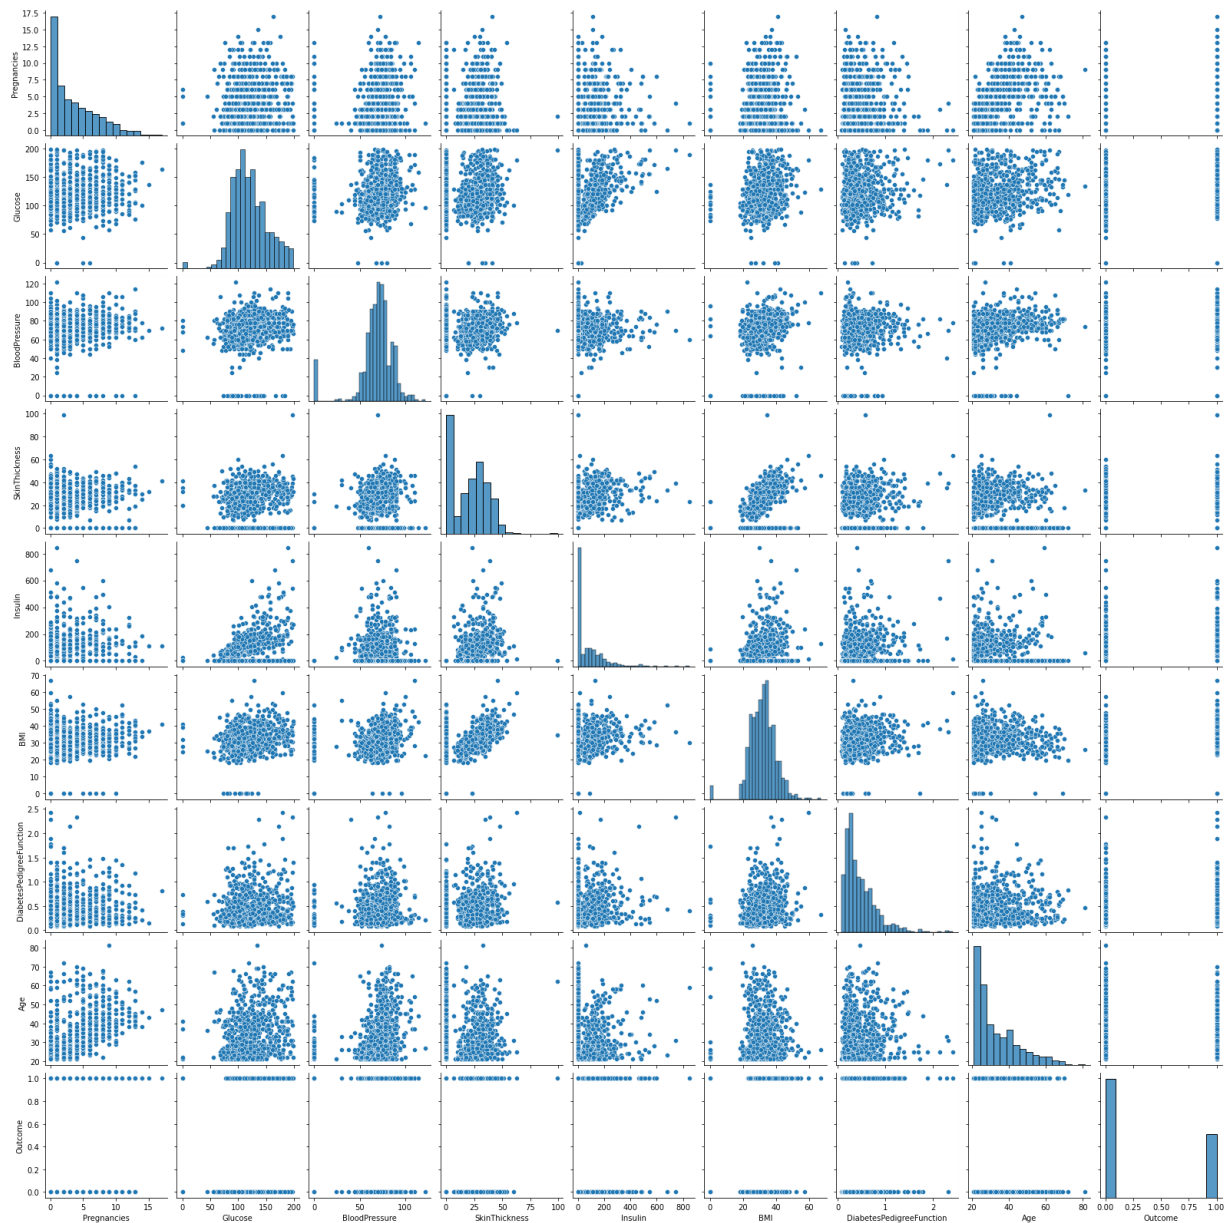
```
Out[10]: Pregnancies      0
         Glucose          5
         BloodPressure    35
         SkinThickness    227
         Insulin          374
         BMI              11
         DiabetesPedigreeFunction  0
         Age              0
```

Outcome
dtype: int64 0

Let's see the distribution of data points in order to fill the null values.

```
In [11]: sns.pairplot(df)
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x249dc9ccd60>
```



Filling Missing Values

```
In [12]: df_copy.isnull().sum()
```

```
Out[12]: Pregnancies      0
Glucose      5
BloodPressure 35
SkinThickness 227
Insulin      374
BMI          11
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

```
In [13]: #Function to find median
```

```
def median_imp(var):
    med_df = df_copy[df_copy[var].notnull()]
    med_df = med_df[[var, 'Outcome']].groupby(['Outcome'])[var].median().reset_index()
    return med_df
```

In [14]: median_imp('Glucose')

Out[14]:

	Outcome	Glucose
0	0	107.0
1	1	140.0

In [15]:

```
df_copy.loc[(df_copy['Outcome']==0) & df_copy['Glucose'].isnull(), 'Glucose'] = 107
df_copy.loc[(df_copy['Outcome']==1) & df_copy['Glucose'].isnull(), 'Glucose'] = 140
```

In [16]: median_imp('BloodPressure')

Out[16]:

	Outcome	BloodPressure
0	0	70.0
1	1	74.5

In [17]:

```
df_copy.loc[(df_copy['Outcome']==0) & df_copy['BloodPressure'].isnull(), 'BloodPressure'] = 70
df_copy.loc[(df_copy['Outcome']==1) & df_copy['BloodPressure'].isnull(), 'BloodPressure'] = 74.5
```

In [18]: median_imp('SkinThickness')

Out[18]:

	Outcome	SkinThickness
0	0	27.0
1	1	32.0

In [19]:

```
df_copy.loc[(df_copy['Outcome']==0) & df_copy['SkinThickness'].isnull(), 'SkinThickness'] = 27
df_copy.loc[(df_copy['Outcome']==1) & df_copy['SkinThickness'].isnull(), 'SkinThickness'] = 32
```

In [20]: median_imp('Insulin')

Out[20]:

	Outcome	Insulin
0	0	102.5
1	1	169.5

In [21]:

```
df_copy.loc[(df_copy['Outcome']==0) & df_copy['Insulin'].isnull(), 'Insulin'] = 102.5
df_copy.loc[(df_copy['Outcome']==1) & df_copy['Insulin'].isnull(), 'Insulin'] = 169.5
```

In [22]: median_imp('BMI')

Out[22]:

	Outcome	BMI
0	0	30.1
1	1	34.3

In [23]:

```
df_copy.loc[(df_copy['Outcome']==0) & df_copy['BMI'].isnull(), 'BMI'] = 30.1
df_copy.loc[(df_copy['Outcome']==1) & df_copy['BMI'].isnull(), 'BMI'] = 34.3
```

```
In [24]: df_copy.isnull().sum()
```

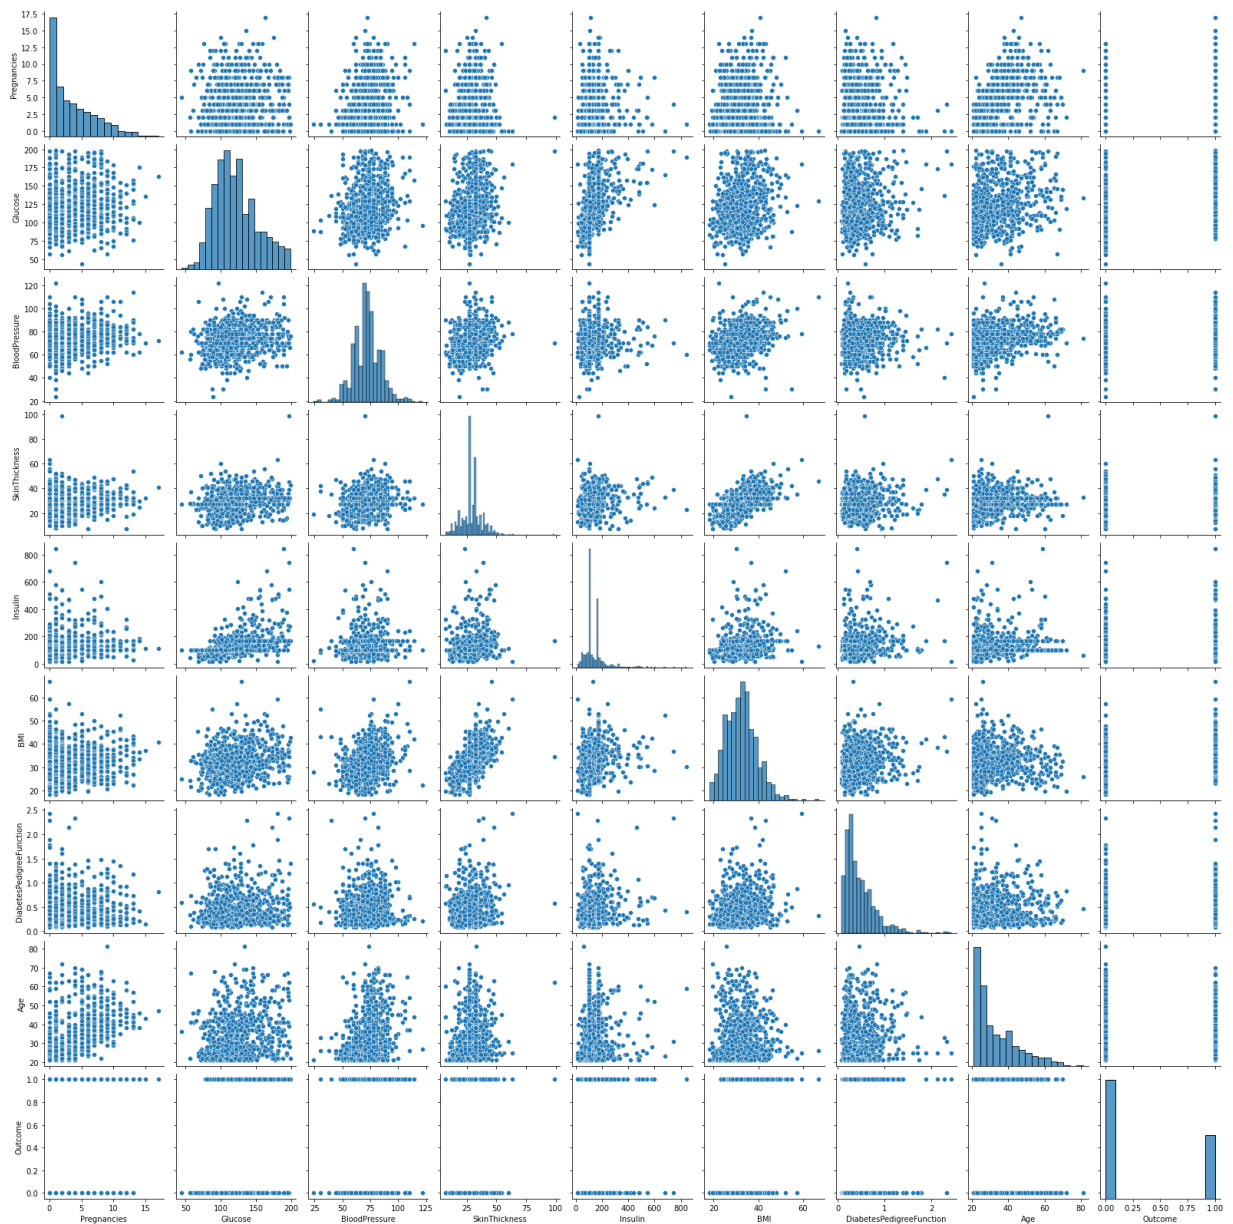
```
Out[24]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

Now our dataset is free from any null values so we can proceed further

Pair Plot after handling missing values

```
In [25]: sns.pairplot(df_copy)
```

```
Out[25]: <seaborn.axisgrid.PairGrid at 0x249dc9aeeb0>
```



Count of types of columns in dataset

```
In [27]: int_dtype = df.select_dtypes(include=['int64']).columns
float_dtype = df.select_dtypes(include=['float64']).columns
obj_dtype = df.select_dtypes(include=['object']).columns
```

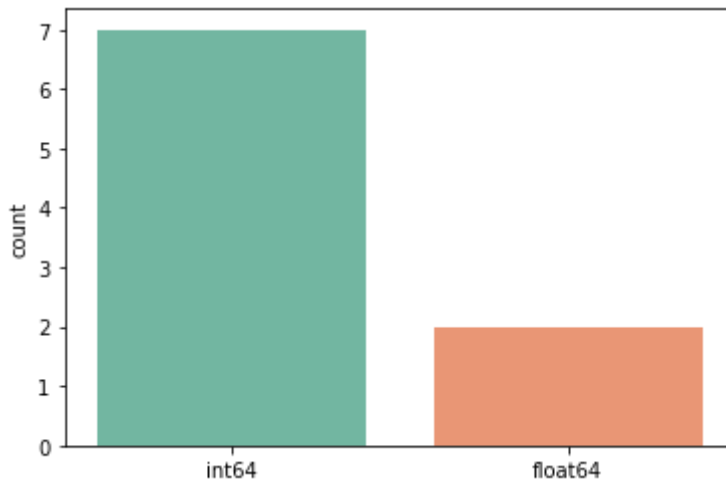


```
In [28]: print('No of integer columns in dataframe is :',len(int_dtype))
        print('No of flaot columns in dataframe is :',len(float_dtype))
        print('No of object columns in dataframe is :',len(obj_dtype))
```

```
No of integer columns in dataframe is : 7
No of flaot columns in dataframe is : 2
No of object columns in dataframe is : 0
```

```
In [29]: sns.countplot(x = df.dtypes.map(str),palette='Set2')
```

```
Out[29]: <AxesSubplot:ylabel='count'>
```



Count of diabetic and healty people in dataset

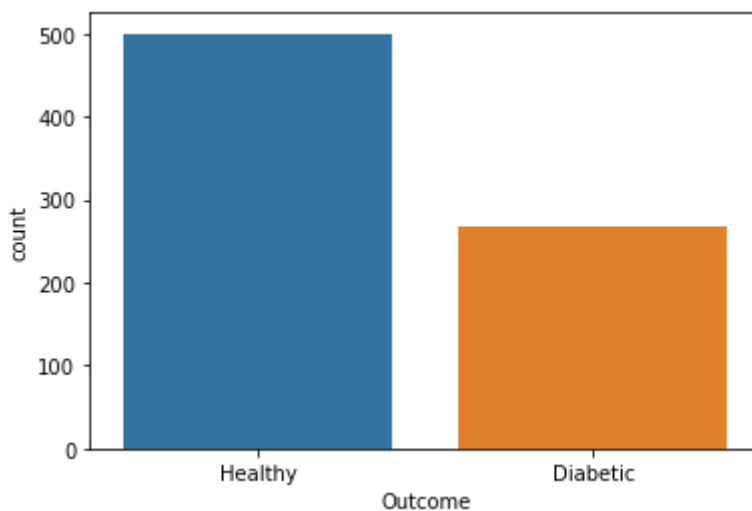
```
In [30]: df.Outcome.value_counts()
```

```
Out[30]: 0    500
        1    268
        Name: Outcome, dtype: int64
```

```
In [31]: diab_count = df.Outcome.astype('category').cat.rename_categories(['Healthy', 'Diabeti
```

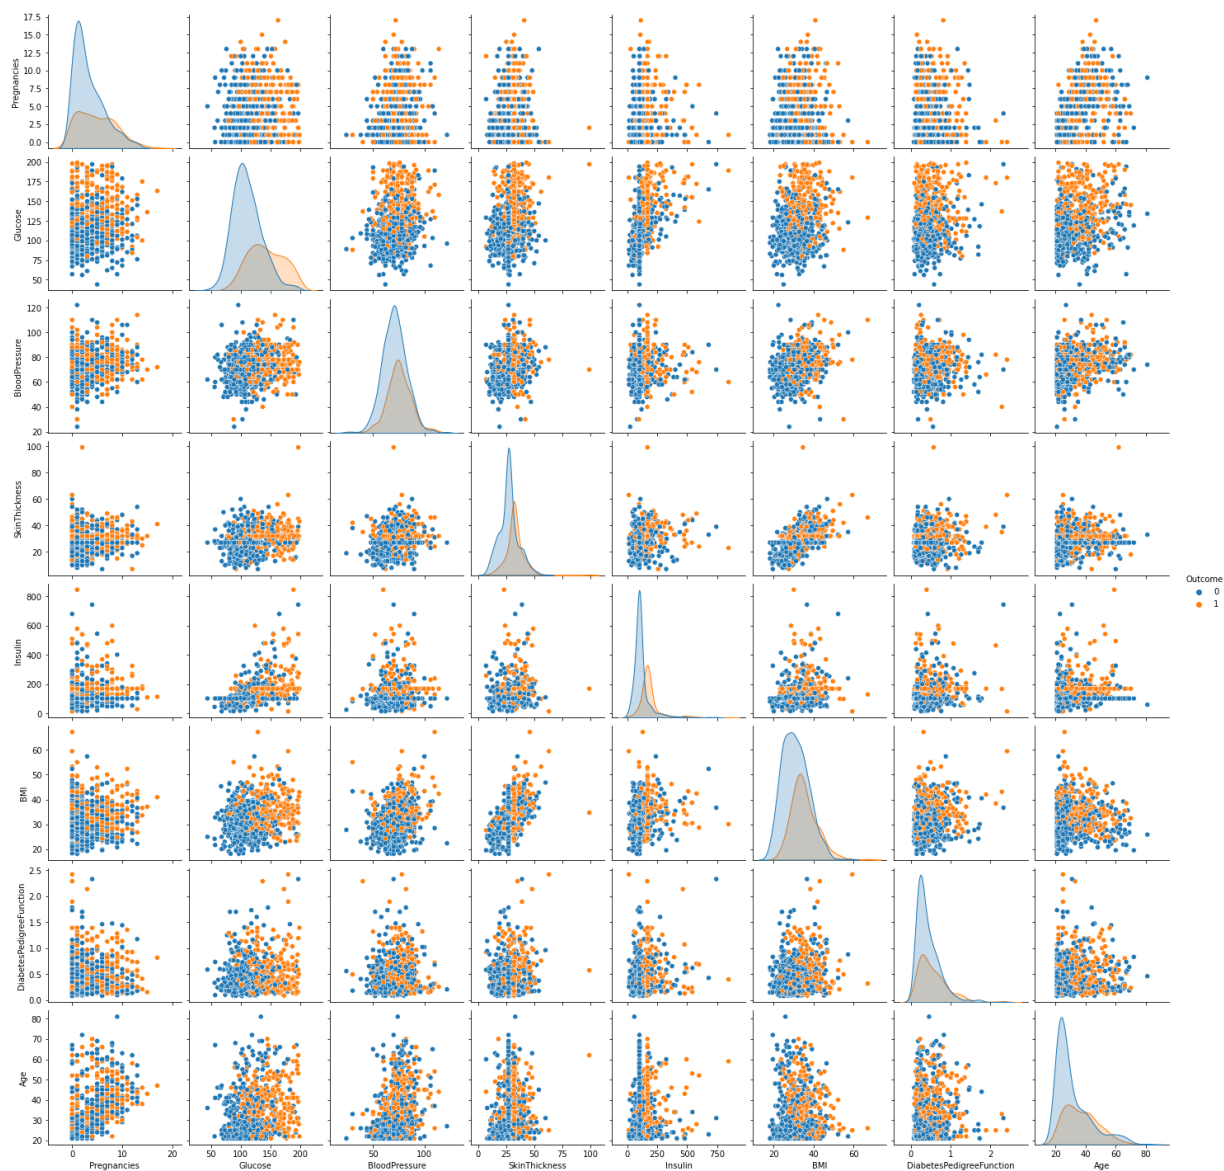
```
In [32]: sns.countplot(x= diab_count)
```

```
Out[32]: <AxesSubplot:xlabel='Outcome', ylabel='count'>
```



```
In [33]: sns.pairplot(df_copy,hue = 'Outcome')
```


```
Out[33]: <seaborn.axisgrid.PairGrid at 0x249e5ecd070>
```

Heatmap of Original Dataset

```
In [34]: plt.figure(figsize=(12,10))
sns.heatmap(df.corr(),annot=True,cmap = 'viridis')
```


```
Out[34]: <AxesSubplot:>
```



Heatmap of Clean Data

```
In [35]: plt.figure(figsize=(12,10))
sns.heatmap(df_copy.corr(),annot=True,cmap = 'viridis')

Out[35]: <AxesSubplot:>
```




From the above heatmap we see a bit of correlation between some columns i.e.

- Age and Pregnancies = 0.54
- Glucose and insulin = 0.49
- SkinThickness and BMI = 0.57


Let's create some scatter plots for above mentioned column pairs to understand the relationship among the top correlation values:

```
In [92]: def sctr_plot(var1,var2):
          sns.scatterplot(x = var1,y = var2, data = df_copy,hue = 'Outcome',marker = 'D')
```


```
In [94]: sctr_plot('Age','Pregnancies')
```



```
In [95]: sctr_plot('Glucose','Insulin')
```



```
In [96]: sctr_plot('SkinThickness','BMI')
```



Data Split for training and testing

```
In [36]: X = df_copy.drop('Outcome',axis=1)
         y=df_copy.Outcome
```

```
In [37]: X.shape,y.shape
```

```
Out[37]: ((768, 8), (768,))
```

```
In [38]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_stat
```

```
In [39]: X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
Out[39]: ((537, 8), (231, 8), (537,), (231,))
```

Standardization

To bring the whole data at a same scale we'll perform standardization.

```
In [40]: sc = StandardScaler()
```

```
In [41]: X_train_scaled = sc.fit_transform(X_train)
```

```
In [42]: X_test_scaled = sc.transform(X_test)
```

1. LogisticRegression Model

```
In [43]: Log_model = LogisticRegression(max_iter=10000)
```

```
In [44]: Log_model.fit(X_train_scaled,y_train)
```

```
Out[44]: LogisticRegression(max_iter=10000)
```

```
In [45]: log_pred = Log_model.predict(X_test_scaled)
```

```
In [46]: print(classification_report(y_test,log_pred))
```

	precision	recall	f1-score	support
0	0.79	0.84	0.81	150
1	0.66	0.58	0.62	81
accuracy			0.75	231
macro avg	0.72	0.71	0.72	231
weighted avg	0.74	0.75	0.74	231

```
In [47]: print(confusion_matrix(y_test,log_pred))
print('\n','Accuracy - ',accuracy_score(y_test,log_pred))
```

```
[[126  24]
 [ 34  47]]
```

```
Accuracy - 0.7489177489177489
```

2. RandomForest Classifier

```
In [48]: rfc = RandomForestClassifier()
```

```
In [49]: n_estimators = [75,100,125,150,200]
```

```
In [50]: max_features = [4,5,6,7,8]
```

```
In [51]: bootstrap=[True,False]
```

```
In [52]: oob_score=[True,False]
```

```
In [53]: param_grid = {'n_estimators':n_estimators,
                       'max_features':max_features,
                       'bootstrap':bootstrap,
                       'oob_score':oob_score}
```

```
In [54]: grid = GridSearchCV(rfc,param_grid)
```

```
In [55]: grid.fit(X_train_scaled,y_train)
```

```
Out[55]: GridSearchCV(estimator=RandomForestClassifier(),
                      param_grid={'bootstrap': [True, False],
                                   'max_features': [4, 5, 6, 7, 8],
                                   'n_estimators': [75, 100, 125, 150, 200],
                                   'oob_score': [True, False]})
```

```
In [56]: rfc_pred = grid.predict(X_test_scaled)
```

```
In [57]: print(classification_report(y_test,rfc_pred))
```

	precision	recall	f1-score	support
0	0.90	0.93	0.92	150
1	0.87	0.80	0.83	81
accuracy			0.89	231
macro avg	0.88	0.87	0.87	231
weighted avg	0.89	0.89	0.89	231

```
In [58]: print(confusion_matrix(y_test,log_pred))
print('\n','Accuracy - ',accuracy_score(y_test,rfc_pred))
```

```
[[126  24]
 [ 34  47]]
```

```
Accuracy - 0.8874458874458875
```

```
In [59]: grid.best_params_
```

```
Out[59]: {'bootstrap': True, 'max_features': 5, 'n_estimators': 125, 'oob_score': True}
```

3. Support Vector Machine

```
In [60]: svc = SVC()
```

```
In [61]: param_grid = {'C':[0.01,0.1,1,10],'kernel':['linear', 'poly', 'rbf', 'sigmoid']}
```

```
In [62]: grid = GridSearchCV(svc,param_grid)
```

```
In [63]: grid.fit(X_train_scaled,y_train)
```

```
Out[63]: GridSearchCV(estimator=SVC(),
                      param_grid={'C': [0.01, 0.1, 1, 10],
                                   'kernel': ['linear', 'poly', 'rbf', 'sigmoid']})
```

```
In [64]: svc_pred = grid.predict(X_test_scaled)
```

```
In [65]: print(classification_report(y_test,svc_pred))
```

	precision	recall	f1-score	support
0	0.85	0.87	0.86	150
1	0.75	0.70	0.73	81
accuracy			0.81	231
macro avg	0.80	0.79	0.79	231
weighted avg	0.81	0.81	0.81	231

```
In [66]: print(confusion_matrix(y_test,log_pred))
print('\n','Accuracy - ',accuracy_score(y_test,svc_pred))

[[126  24]
 [ 34  47]]

Accuracy - 0.8138528138528138
```

```
In [67]: grid.best_params_
```


```
Out[67]: {'C': 1, 'kernel': 'rbf'}
```

4. K Nearest Neighbour

```
In [68]: train_score = []
test_score = []
for i in range(1,20):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train_scaled,y_train)
    # knn_pred = knn.predict(X_test)
    train_score.append(knn.score(X_train_scaled,y_train))
    test_score.append(knn.score(X_test_scaled,y_test))

sns.lineplot(x = range(1,20),y = train_score,marker='o')
sns.lineplot(x = range(1,20),y = test_score,marker = 'o')
```

```
Out[68]: <AxesSubplot:>
```



```
In [69]: acc_score = []
for i in range(1,10):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train_scaled,y_train)
    knn_pred = knn.predict(X_test_scaled)


    acc_score.append(accuracy_score(y_test,knn_pred))
print(max(acc_score))
```



```
sns.lineplot(x = range(1,10),y = acc_score,marker='o')
```

```
0.8398268398268398
```

```
Out[69]: <AxesSubplot:>
```



From above results this could be concluded that $n=7$ gives the best results so we'll take $n_neighbors = 7$ for final model

```
In [70]: final_knn_model = KNeighborsClassifier(n_neighbors=7)
```

```
In [71]: final_knn_model.fit(X_train_scaled,y_train)
```

```
Out[71]: KNeighborsClassifier(n_neighbors=7)
```

```
In [72]: knn_pred = final_knn_model.predict(X_test_scaled)
```

```
In [73]: print(accuracy_score(y_test,knn_pred))
```

```
0.8398268398268398
```

```
In [74]: print(classification_report(y_test,knn_pred))
```

	precision	recall	f1-score	support
0	0.87	0.89	0.88	150
1	0.78	0.75	0.77	81
accuracy			0.84	231
macro avg	0.83	0.82	0.82	231
weighted avg	0.84	0.84	0.84	231

```
In [75]: print(confusion_matrix(y_test,knn_pred))
```

```
[[133  17]
 [ 20  61]]
```

```
In [ ]:
```

5. Decision Tree

```
In [76]: dt = DecisionTreeClassifier(random_state=42)
```

```
In [77]: param_grid = {'criterion' : ["gini", "entropy"],
                      'min_samples_split' : [2,3,4,5],
```

```
'max_features' : [4,5,6,7,8]
}
```

```
In [78]: grid_dt = GridSearchCV(dt,param_grid)
```

```
In [79]: grid_dt.fit(X_train_scaled,y_train)
```

```
Out[79]: GridSearchCV(estimator=DecisionTreeClassifier(random_state=42),
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_features': [4, 5, 6, 7, 8],
                                   'min_samples_split': [2, 3, 4, 5]})
```

```
In [80]: grid_dt.best_params_
```

```
Out[80]: {'criterion': 'entropy', 'max_features': 6, 'min_samples_split': 3}
```

```
In [81]: dt_pred = grid_dt.predict(X_test_scaled)
```

```
In [82]: print(accuracy_score(y_test,dt_pred))
```

```
0.8614718614718615
```

```
In [83]: print(classification_report(y_test,dt_pred))
```

	precision	recall	f1-score	support
0	0.87	0.92	0.90	150
1	0.84	0.75	0.79	81
accuracy			0.86	231
macro avg	0.85	0.84	0.84	231
weighted avg	0.86	0.86	0.86	231

```
In [104... dt_cm = confusion_matrix(y_test,dt_pred)
```

```
In [105... print(confusion_matrix(y_test,dt_pred))
```

```
[[138 12]
 [ 20 61]]
```

```
In [85]: grid_dt.best_params_
```

```
Out[85]: {'criterion': 'entropy', 'max_features': 6, 'min_samples_split': 3}
```

```
In [111...
```

In a Nutshell

Accuracy , Sensitivity and Specificity

```
In [124... models = [
{
    'label': 'Logistic Regression',
    'model': LogisticRegression(),
},
{
    'label': 'KNeighbors Classifier',
    'model': KNeighborsClassifier(n_neighbors=7),
},
{
```

```

    'label' : 'Support Vector Classifier',
    'model' : SVC(C= 1, kernel='rbf',probability=True),
},
{
    'label' : 'Decision Tress',
    'model' : DecisionTreeClassifier(random_state=42,criterion= 'entropy', max_featu
},
{
    'label' : 'Random Forest Classifier',
    'model' : RandomForestClassifier(bootstrap= True, max_features= 6, n_estimators=
},
]

```

In [136...

```

accu = []
model_name= []
sensitivity = []
specificity = []
for m in models:
    model1 = m['model']
    model1.fit(X_train_scaled, y_train) # train the model
    pred = model1.predict(X_test_scaled) # predict the test data
    cm = confusion_matrix(y_test,pred)

    accu.append(accuracy_score(y_test,pred))
    model_name.append(m['label'])
    # sensitivity.append(cm[0,0]/(cm[0,0]+cm[0,1]))
    # specificity.append(cm[1,1]/(cm[1,0]+cm[1,1]))

    models_accuracy= pd.DataFrame(data=accu,index = model_name,columns=['Accuracy Sc
models_accuracy

```

Out[136...

	Accuracy Score
Logistic Regression	0.748918
KNeighbors Classifier	0.839827
Support Vector Classifier	0.813853
Decision Tress	0.861472
Random Forest Classifier	0.883117

In [182...

```

accu = []
model_name= []
sensitivity = []
specificity = []
for m in models:
    model1 = m['model']
    model1.fit(X_train_scaled, y_train) # train the model
    pred = model1.predict(X_test_scaled) # predict the test data
    cm = confusion_matrix(y_test,pred)

    accu.append(accuracy_score(y_test,pred))
    model_name.append(m['label'])

    sensitivity.append(cm[0,0]/(cm[0,0]+cm[0,1]))
    specificity.append(cm[1,1]/(cm[1,0]+cm[1,1]))

    models_accu_sen_sp= pd.DataFrame(data=(accu,sensitivity,specificity),index = ['A
                                columns=[model_name]).T
models_accu_sen_sp

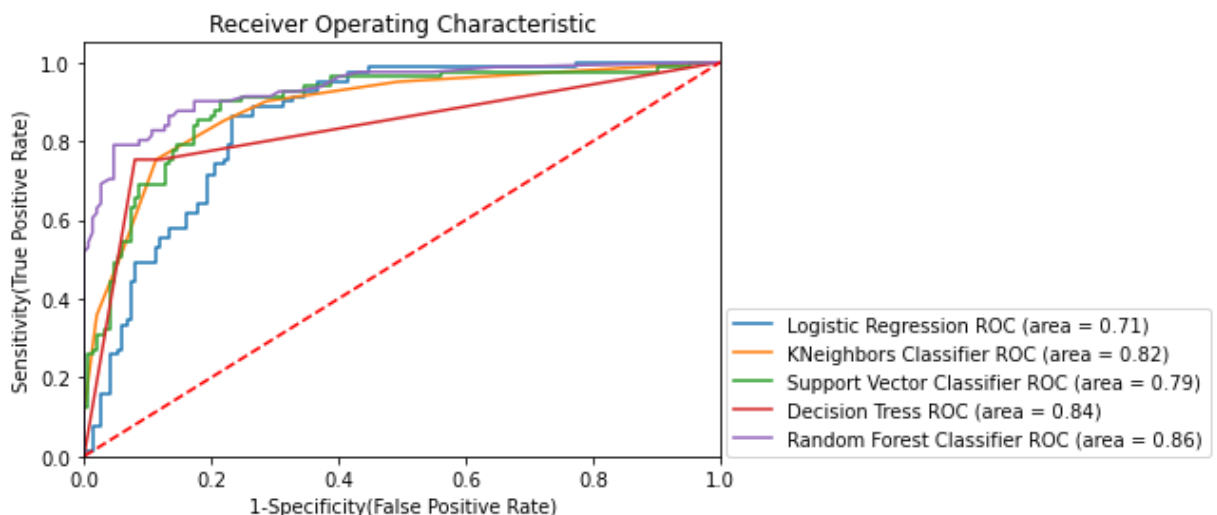
```

Out[182...

	Accuracy	Sensitivity	Specificity
Logistic Regression	0.748918	0.840000	0.580247
KNeighbors Classifier	0.839827	0.886667	0.753086
Support Vector Classifier	0.813853	0.873333	0.703704
Decision Tress	0.861472	0.920000	0.753086
Random Forest Classifier	0.874459	0.920000	0.790123

Combined ROC Curve for all the models

```
In [88]: # Below for Loop iterates through your models List
for m in models:
    model = m['model'] # select the model
    model.fit(X_train_scaled, y_train) # train the model
    y_pred=model.predict(X_test_scaled) # predict the test data
    # Compute False positive rate, and True positive rate
    fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test_scaled)[: ,1])
    # Calculate Area under the curve to display on the plot
    auc = roc_auc_score(y_test,model.predict(X_test_scaled))
    # Now, plot the computed values
    plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % (m['label'], auc))
    # Custom settings for the plot
    plt.plot([0, 1], [0, 1], 'r--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('1-Specificity(False Positive Rate)')
    plt.ylabel('Sensitivity(True Positive Rate)')
    plt.title('Receiver Operating Characteristic')
    plt.legend(loc=(1.01,0))
    plt.show() # Display
```




So the above curve gives the max area under the curve for Random Forest Classifier, hence this could be the best algorithm for the problem.

Tableau Visualisation

```
In [199... image1 = PILImage.open('C:/Users/vipul/ML project 1/Tableau Dashboard/Dashboard 1.png')
image1
```


Out[199...



In [195...

```
image2 = PILImage.open('C:/Users/vipul/ML project 1/Tableau Dashboard/Dashboard 4.png')
image2
```


Out[195...



In [196...

```
image3 = PILImage.open('C:/Users/vipul/ML project 1/Tableau Dashboard/Scatter Plot a.png')
image3
```

Out[196...




In [197...

```
image4 = PILImage.open('C:/Users/vipul/ML project 1/Tableau Dashboard/Scatter Plot.png')
image4
```

Out[197...

Scatter Plot



In []: